

dataExploration

```
// val github = img()  
// github.url(new java.net.URL("https://assets-cdn.github.com/images/modules/logos_page/GitHub-Mark.png"))  
// github
```

Took: 3 seconds 49 milliseconds, at 2017-5-5 1:8

Import libraries

```
import org.apache.spark._  
import org.apache.spark.SparkContext._  
import org.apache.spark.rdd._  
import org.apache.spark.sql.SparkSession  
  
import org.apache.spark.sql.types.{DataType}  
import org.apache.spark.sql.types.IntegerType  
import org.apache.spark.sql.types.FloatType  
import org.apache.spark.sql.types.LongType  
  
val spark = SparkSession.builder().getOrCreate()  
import spark.implicits._
```

Took: 4 seconds 648 milliseconds, at 2017-5-5 1:8

Import train set

```
val trainRaw = spark.read
    .format("csv")
    .option("header", "true")
    .option("mode", "DROPMALFORMED")
    .csv("/Users/Albert/Data Science/Kaggle/Facebook/train.csv")

trainRaw.show()
```

Took: 7 seconds 956 milliseconds, at 2017-5-5 1:8

Count the number of imported rows and check the DataFrame schema

```
trainRaw.count()
```

29118021

Took: 41 seconds 733 milliseconds, at 2017-5-5 1:9

```
trainRaw.printSchema()
```

Took: 2 seconds 449 milliseconds, at 2017-5-5 1:9

Recast columns to appropriate data types before summarising

```
val train = trainRaw.select(trainRaw("row_id").cast(IntegerType).as("row_id"),
                           trainRaw("x").cast(FloatType).as("x"),
                           trainRaw("y").cast(FloatType).as("y"),
                           trainRaw("accuracy").cast(IntegerType).as("accuracy"),
                           trainRaw("time").cast(IntegerType).as("time"),
                           round(trainRaw("time").cast(IntegerType), -3).as("timeRounded"),
                           trainRaw("place_id").cast(LongType).as("place_id")
                          )

train.printSchema()
```

Took: 2 seconds 894 milliseconds, at 2017-5-5 1:9


```
val trainAccHist = train.groupBy("accuracy").count()
val trainTimeHist = train.groupBy("time").count()
val trainTimeRoundedHist = train.groupBy("timeRounded").count()
val trainPlaceHist = train.groupBy("place_id").count()

val trainXAvgAcc = train.groupBy("x").agg(mean(train("accuracy")))
val trainYAvgAcc = train.groupBy("y").agg(mean(train("accuracy")))
```


Took: 2 seconds 46 milliseconds, at 2017-5-5 1:9

Plot charts

```
trainAccHist.collect()
```




1025 entries total (Warning: randomly sampled 1000 entries)




Took: 51 seconds 8

```
trainTimeHist.collect()
```

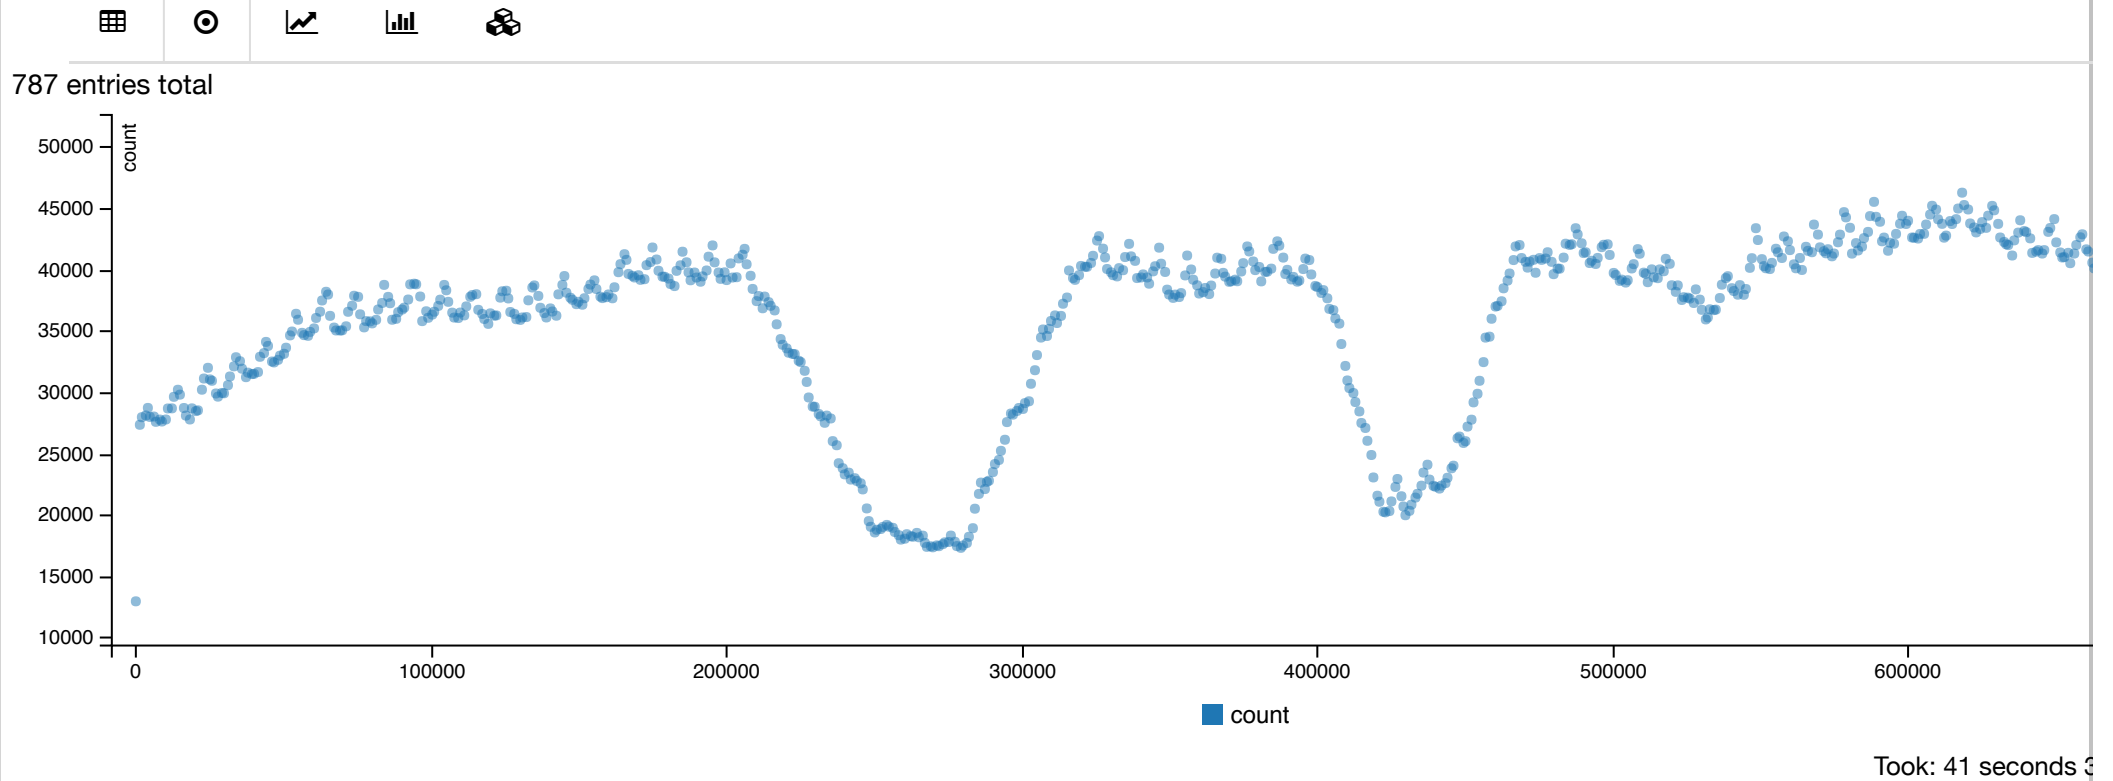


786239 entries total (Warning: randomly sampled 1000 entries)



Took: 52 seconds 5


```
trainTimeRoundedHist.collect()
```



```
trainPlaceHist.collect()
```




108390 entries total (Warning: randomly sampled 1000 entries)



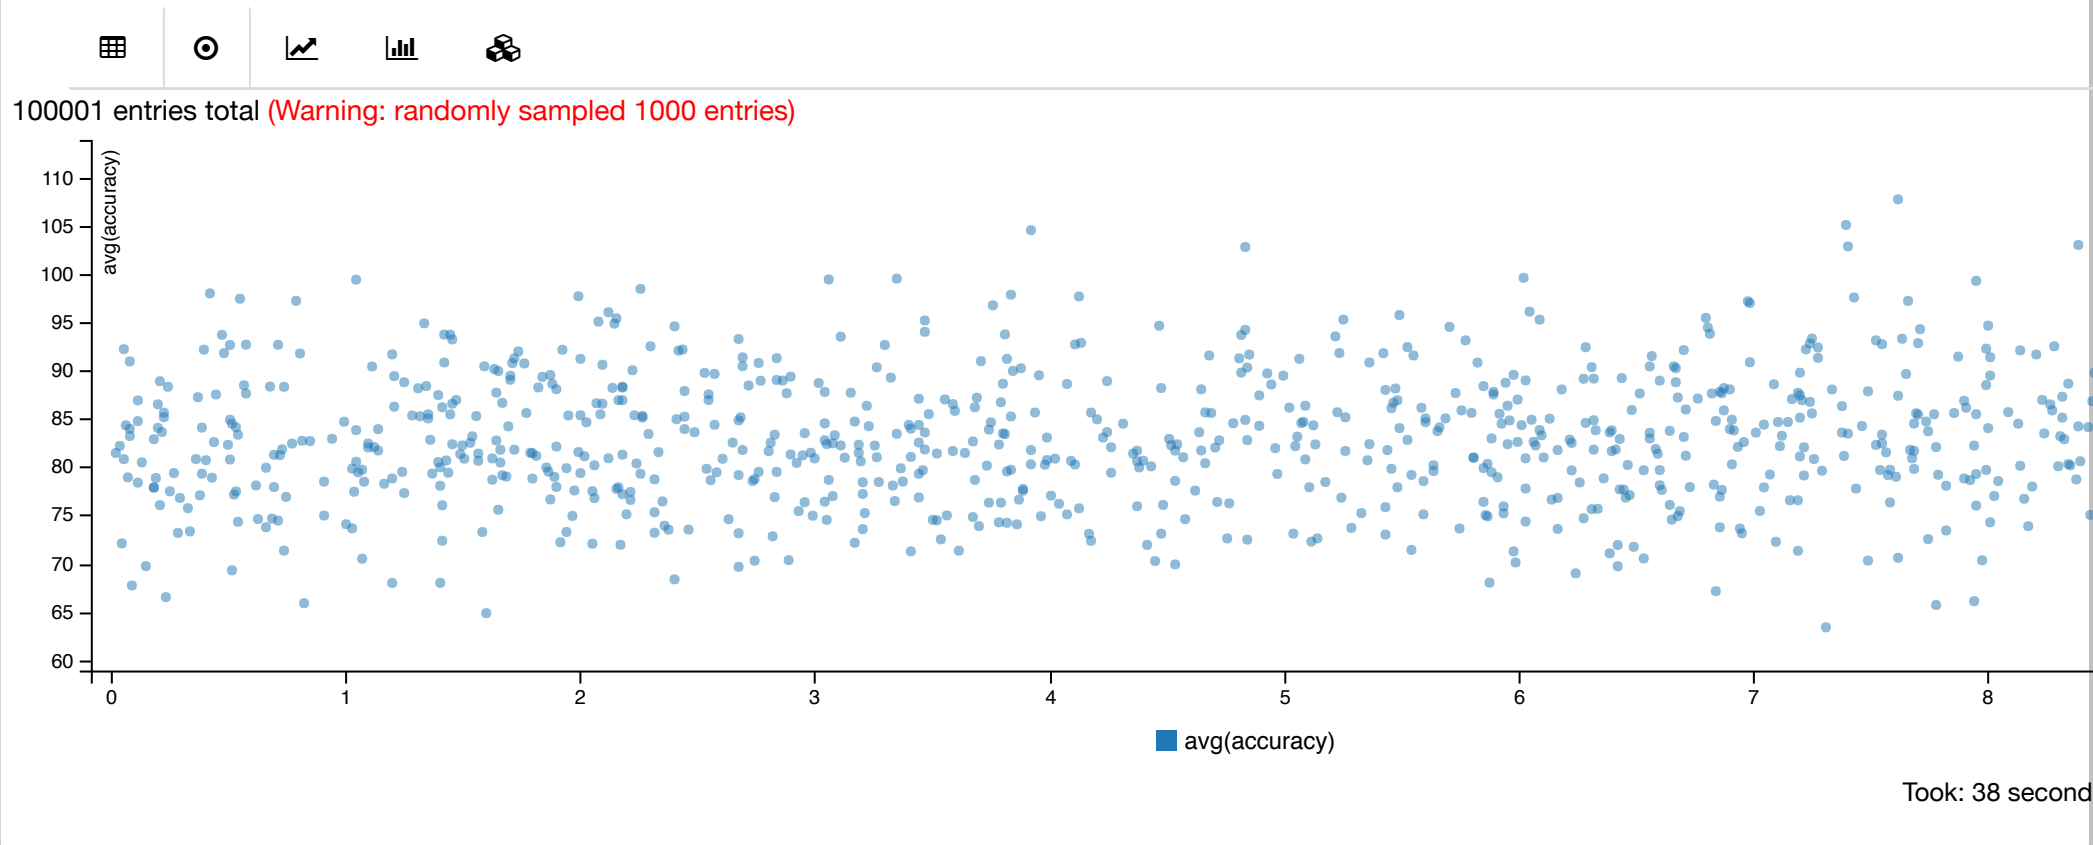
```
trainXAvgAcc.collect()
```



100001 entries total (Warning: randomly sampled 1000 entries)




```
trainYAvgAcc.collect()
```



Build: | **buildTime**-Mon Jan 30 17:32:36 UTC 2017 | **formattedShaVersion**-0.7.0-c955e71d0204599035f603109527e679aa3bd570 | **sbtVersion**-0.13.8 | **scalaVersion**-2.11.8 | **sparkNotebookVersion**-0.7.0 | **hadoopVersion**-2.7.3 | **jets3tVersion**-0.7.1 | **jlineDef**-(jline,2.12) | **sparkVersion**-2.1.0 | **withHive**-true |.